

Python basics

NLTK

Texts

Lists

Distributions

Control structures

Nested Blocks

New data

POS Tagging

Basic tagging

Tagged corpora

Automatic tagging

The Natural Language Toolkit (NLTK)

Markus Dickinson

Dept. of Linguistics, Indiana University
Catapult Workshop Series; March 8, 2013

Natural Language Toolkit (NLTK) is:

Open source Python modules, linguistic data and documentation for research and development in natural language processing and text analytics, with distributions for Windows, Mac OSX and Linux.

<http://www.nltk.org/>

Today, we'll look at:

- ▶ Some basic functionality for working with text files
 - ▶ <http://nltk.org/book/ch01.html>
 - ▶ <http://nltk.org/book/ch03.html>
- ▶ One example of an NLP process, POS tagging
 - ▶ <http://nltk.org/book/ch05.html>

Python basics

NLTK

Texts

Lists

Distributions

Control structures

Nested Blocks

New data

POS Tagging

Basic tagging

Tagged corpora

Automatic tagging

Where we're going

NLTK is a package written in the programming language Python, providing a lot of tools for working with text data

Goals: By the end of today, you should be:

- ▶ Familiar enough with Python to work with NLTK
- ▶ Familiar with NLTK, so as to be able to:
 - ▶ Use their pre-installed data files
 - ▶ Import your own text data files
 - ▶ Employ basic data manipulation
 - ▶ Part-of-speech (POS) tag your data
- ▶ Comfortable with NLTK, so as to be able to teach yourself more on other NLP applications, e.g.,
 - ▶ Classification
 - ▶ Parsing, Chunking, & Grammar Writing
 - ▶ Propositional Semantics & Logic

Python basics

NLTK

Texts

Lists

Distributions

Control structures

Nested Blocks

New data

POS Tagging

Basic tagging

Tagged corpora

Automatic tagging

NLTK is based on Python

- ▶ We will assume Python 2.7 for now, but Python 3 is the way to go for the future ...
 - ▶ Python 3 is Unicode all the way through, allowing for easy handling of various languages
- ▶ Python is a full programming language

Python has two modes:

- ▶ **Interactive** → our focus today
- ▶ File-based

There are a lot of good Python resources out there:

- ▶ The main Python tutorial:
<http://www.python.org/doc/current/tut/>
- ▶ Code Academy:
<http://www.codecademy.com/tracks/python>
- ▶ etc.

Python basics

NLTK

Texts

- Lists
- Distributions
- Control structures
- Nested Blocks
- New data

POS Tagging

- Basic tagging
- Tagged corpora
- Automatic tagging

To start, type `python` in a terminal or command prompt

- ▶ Better yet might be to use the Interactive DeveLopment Environment (IDLE)

```
> python
```

```
Python 2.7.2 (default, Jun 20 2012, 16:23:33)
```

```
[GCC 4.2.1 Compatible Apple Clang 4.0 (tags/Apple/clang-418.0.60)]
```

```
Type "help", "copyright", "credits" or "license" for more informati
```

```
>>>
```

Some uses of numbers:

```
>>> 2+2
4
>>> 3/2.
1.5
```

Some uses of strings:

- ▶ single quotes: 'string'
- ▶ double quotes: "string"
- ▶ There are string characters with special meaning: e.g.,
\\n (newline) and \\t (tab)

String indices & slices

You can use slices to get a part of a string

```
>>> s = "happy"
>>> len(s)    # use the len function
5
>>> s[3]      # indexed from 0, so 4th character
'p'
>>> s[1:3]    # characters 1 and 2
'ap'
>>> s[:3]     # first 3 characters
'hap'
>>> s[3:]     # everything except first 3 characters
'py'
>>> s[-4]     # 4th character from the back
'a'
```

Python basics

NLTK

Texts

Lists

Distributions

Control structures

Nested Blocks

New data

POS Tagging

Basic tagging

Tagged corpora

Automatic tagging

Definition

A variable is a name that refers to some value (could be a number, a string, a list etc.)

1. Store the value 42 in a variable named *foo*

```
foo = 42
```

2. Store the value of *foo*+10 in a variable named *bar*

```
bar = foo + 10
```


Installing NLTK is pretty straightforward:

- ▶ <http://nltk.org/install.html>
- ▶ I recommend installing Numpy, but that can sometimes present challenges

Download the materials from the NLTK book:

```
>>> import nltk
>>> nltk.download()
...
Downloader> d book
...
```

This command gives us various texts to work with, which we need to load:

```
>>> from nltk.book import *
```

We now have texts available:

```
>>> text1  
<Text: Moby Dick by Herman Melville 1851>
```

Methods which do some basic analysis:

- ▶ concordance
text1.concordance("monstrous")
- ▶ similar
text1.similar("monstrous")
- ▶ common_contexts
text2.common_contexts(["monstrous", "very"])

NLTK treats texts as lists of words (more on lists in a bit):

Here are the first 20 words of *Moby Dick*

```
>>> text1[:20]  
['[', 'Moby', 'Dick', 'by', 'Herman', 'Melville',  
'1851', ']', 'ETYMOLOGY', '.', '(', 'Supplied',  
'by', 'a', 'Late', 'Consumptive', 'Usher', 'to',  
'a', 'Grammar']
```

Because it's Python-based, it's easy to create functions to analyze the texts

```
>>> from __future__ import division
>>> def lexical_diversity(text):
...     return len(text) / len(set(text))
...
>>> lexical_diversity(text1)
13.502044830977896
>>> lexical_diversity(text2)
20.719449729255086
```

Note: `set()` converts a list to a set

- ▶ If you're not familiar with sets, check 'em out! (<http://docs.python.org/2/tutorial/datastructures.html#sets>)

Python basics

NLTK

Texts

Lists

Distributions

Control structures

Nested Blocks

New data

POS Tagging

Basic tagging

Tagged corpora

Automatic tagging

Let's detour back to Python to talk about lists:

- ▶ Lists are containers for more than one element
 - ▶ example: `employee = ['Markus', 'Dickinson', 'assistant prof', 'MM317']`
 - ▶ empty list: `[]`
- ▶ Each element in the sequence is assigned a position number, an **index** (starting from 0)
 - ▶ example: `employee[1]`
- ▶ **Indexing**: accessing elements in a list
 - ▶ `greeting = ['hi', 'there', 'partner']`
`greeting[2]`
- ▶ Adding lists:
`long_greeting = greeting + ['how', 'are', 'you']`

Python basics

NLTK

Texts

Lists

Distributions

Control structures

Nested Blocks

New data

POS Tagging

Basic tagging

Tagged corpora

Automatic tagging

- ▶ accessing parts of segments is called **slicing**
 - ▶ example:
`long_greeting[3:6]`
 - ▶ the slice starts at the first index and goes up to the second (non-inclusive)!
- ▶ going all the way to the end:
`long_greeting[3:]`
- ▶ starting at the beginning:
`long_greeting[:3]`
- ▶ steps are given as optional third number:
`long_greeting[1:6:2]`

- ▶ membership:
`employee = ['Markus, Dickinson', 'assistant
prof', 'MM317']`
`'MM317' in employee`
- ▶ check length:
`len(employee)`
- ▶ add at the end: append
`employee.append('Computational Linguistics')`
- ▶ retrieve from the end: pop
`employee.pop()`
 - ▶ This returns a value!
- ▶ add at the beginning:
`employee.insert(0, 'Linguistics')`
- ▶ retrieve from the beginning:
`employee.pop(0)`

Simple Statistics

Frequency Distributions

NLTK has pre-built packages for creating distributions

```
>>> fdist1 = FreqDist(text1)
>>> fdist1
<FreqDist with 19317 samples and 260819 outcomes>
>>> fdist1['whale']
906
```

You could build your own dictionaries, but some capabilities are quickly calculated with `FreqDist()`:

```
>>> vocabulary1 = fdist1.keys()
>>> vocabulary1[:10]
['', 'the', '.', 'of', 'and', 'a', 'to', ';',
 'in', 'that']
```

Organizing by word length

```
>>> fdist = FreqDist([len(w) for w in text1])
>>> fdist
<FreqDist with 19 samples and 260819 outcomes>
>>> fdist.keys()
[3, 1, 4, 2, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, ...]
>>> fdist.items()
[(3, 50223), (1, 47933), (4, 42345), (2, 38513), ...]
>>> fdist.max()
3
>>> fdist[3]
50223
>>> fdist.freq(3)
0.19255882431878046
```

In the previous slide, what was happening here:

- ▶ `[len(w) for w in text1]`

To answer this, we need to discuss control structures:

- ▶ Conditionals: `if/elif/else`
- ▶ Loops: `for` and `while`
 - ▶ We'll only look at `for` today

If Statement

- ▶ syntax:

```
if <test>:  
    do this
```

- ▶ full program:

```
known_users = [ 'Sandra', 'Markus' ]  
name = raw_input( 'type your name: ' )
```

```
if name in known_users:  
    print 'Hello ' + name
```

- ▶ a test (in the if statement) corresponds to a yes/no question and can be either true or false
- ▶ the following values count as false:

False

None

0

[] (empty list)

{ } (empty dict)

' ' (empty string)

() (empty tuple)

- ▶ everything else counts as true!

Python basics

NLTK

Texts

Lists

Distributions

Control structures

Nested Blocks

New data

POS Tagging

Basic tagging

Tagged corpora

Automatic tagging

Else Statements

- ▶ In case the program needs to do something when the test is false, use the `else:` statement
- ▶ E.g. if a user is not known, add him/her to the list

Example

```
known_users = ['Sandra', 'Markus']  
name = raw_input('type your name: ')  
  
if name in known_users:  
    print 'Hello ' + name + '.'  
    print 'It is nice to have you back.'  
else:  
    known_users.append(name)  
    print 'You have been added to the list.'
```

Python basics

NLTK

Texts

Lists

Distributions

Control structures

Nested Blocks

New data

POS Tagging

Basic tagging

Tagged corpora

Automatic tagging

- ▶ if you want to check the next condition in the else case, there is a shortcut for *else if* called `elif`

Example

```
known_users = ['Sandra', 'Markus']
name = raw_input('type your name: ')

if name in known_users:
    print 'Hello ' + name + ' .'
    print 'It is nice to have you back.'
elif len(name) > 20:
    print 'Your name is too long!'
else:
    known_users.append(name)
    print 'You have been added to the list.'
```

Python basics

NLTK

Texts

Lists

Distributions

Control structures

Nested Blocks

New data

POS Tagging

Basic tagging

Tagged corpora

Automatic tagging

<code>x == y</code>	x equals y
<code>x < y</code>	x is less than y
<code>x > y</code>	x is greater than y
<code>x >= y</code>	x is greater than or equal to y
<code>x <= y</code>	x is less than or equal to y
<code>x != y</code>	x is not equal to y
<code>x is y</code>	x is the same object as y
<code>x is not y</code>	x is not the same object as y
<code>x in y</code>	x is a member of y
<code>x not in y</code>	x is not a member of y

Python basics

NLTK

Texts

Lists

Distributions

Control structures

Nested Blocks

New data

POS Tagging

Basic tagging

Tagged corpora

Automatic tagging

Word comparison tests

<code>s.startswith(t)</code>	test if <code>s</code> starts with <code>t</code>
<code>s.endswith(t)</code>	test if <code>s</code> ends with <code>t</code>
<code>t in s</code>	test if <code>t</code> is contained inside <code>s</code>
<code>s.islower()</code>	test if all cased characters in <code>s</code> are lowercase
<code>s.isupper()</code>	test if all cased characters in <code>s</code> are uppercase
<code>s.isalpha()</code>	test if all characters in <code>s</code> are alphabetic
<code>s.isalnum()</code>	test if all characters in <code>s</code> are alphanumeric
<code>s.isdigit()</code>	test if all characters in <code>s</code> are digits
<code>s.istitle()</code>	test if <code>s</code> is titlecased (all words in <code>s</code> have initial capitals)

Python basics

NLTK

Texts

Lists

Distributions

Control structures

Nested blocks

New data

POS Tagging

Basic tagging

Tagged corpora

Automatic tagging

Iteration

for loops allow us to iterate over each element of a set or sequence

Syntax:

```
for <var> in <set>:  
    do ...  
    do ...
```

Example

```
words = [ 'a', 'rose', 'is', 'a', 'rose', 'is',  
          'a', 'rose']  
  
for w in words:  
    print w
```

List comprehensions

Python has a cool shorthand called **list comprehensions** for creating new lists from old ones:

- ▶ `a = [1,2,3,4,5]`
`b = [x**2 for x in a]`
b is set to [1, 4, 9, 16, 25]

So: `[len(w) for w in text1]` gives a list of word lengths

- ▶ What does this do?

```
sorted([w for w in set(text1)
        if w.endswith('ableness')])
```

Returning to NLTK functions ...

- ▶ Get bigrams from a text (or list):

```
>>> bigrams(text1[:10])  
[(['', 'Moby'), ('Moby', 'Dick'), ('Dick', 'by'),  
 ('by', 'Herman'), ('Herman', 'Melville'),  
 ('Melville', '1851'), ('1851', '']'),  
 (']', 'ETYMOLOGY'), ('ETYMOLOGY', '.')]
```

- ▶ Get the most frequent collocations:

```
>>> text1.collocations()  
Building collocations list  
Sperm Whale; Moby Dick; White Whale; old man;  
Captain Ahab; sperm whale; Right Whale;  
Captain Peleg; New Bedford; Cape Horn; cried Ahab;  
years ago; lower jaw; never mind; Father Mapple; ...
```

Using your own data

Using `.read()`, you can read a text file as a string in Python

- ▶ With a string representation, you can use NLTK's utilities

`raw` is *Crime and Punishment*, from Project Gutenberg

```
>>> raw = open('crime.txt').read()
>>> tokens = nltk.word_tokenize(raw)
>>> tokens[:10]
['The', 'Project', 'Gutenberg', 'EBook', 'of',
 'Crime', 'and', 'Punishment', ',', 'by']
```

- ▶ `open()` opens a file & `read()` converts it to a string

Creating an NLTK text

`nltk.Text()` creates a NLTK text, with all its internal methods available:

```
>>> text = nltk.Text(tokens)
>>> type(text)
<class 'nltk.text.Text'>
>>> text[:10]
['The', 'Project', 'Gutenberg', 'EBook', 'of',
'Crime', 'and', 'Punishment', ',', 'by']
>>> text.collocations()
Building collocations list
Katerina Ivanovna; Pyotr Petrovitch;
Pulcheria Alexandrovna; Avdotya Romanovna;
Marfa Petrovna; Rodion Romanovitch;
Sofya Semyonovna; old woman; Project Gutenberg-tm;
Porfiry Petrovitch; Amalia Ivanovna; great deal; ...
```

Python basics

NLTK

Texts

Lists

Distributions

Control structures

Nested Blocks

New data

POS Tagging

Basic tagging

Tagged corpora

Automatic tagging

Managing corpora in NLTK

There is much more you can do to use your own corpus data in NLTK

- ▶ Some of this involves using Corpus Readers
- ▶ See: <http://nltk.googlecode.com/svn/trunk/doc/howto/corpus.html>

There are options for normalizing words, as well

```
>>> porter = nltk.PorterStemmer()
>>> lancaster = nltk.LancasterStemmer()
>>> [porter.stem(t) for t in tokens]
['DENNI', ':', 'Listen', ',', 'strang',
 'women', 'lie', ...]
>>> [lancaster.stem(t) for t in tokens]
['den', ':', 'list', ',', 'strange',
 'wom', 'lying', ...]
```

We can use NLTK to perform a variety of NLP tasks

- ▶ We will quickly cover the utilities for POS tagging
- ▶ Other modules include:
 - ▶ Classification
 - ▶ Parsing, Chunking, & Grammar Writing
 - ▶ Propositional Semantics & Logic

Segmentation & Tokenization

As we saw, you can use `nltk.word_tokenize()` to break a sentence into tokens

- ▶ `nltk.sent_tokenize` breaks a text into sentences

```
>>> nltk.sent_tokenize("Hello, you fool.  I love\  
... you.  Come join the joyride.")  
['Hello, you fool.', 'I love you.',  
 'Come join the joyride.']
```

Basic NLTK tagging

A very basic way to tag:

```
>>> import nltk
text = nltk.word_tokenize("They refuse to permit us
                           to obtain the refuse permit")
>>> nltk.pos_tag(text)
[('They', 'PRP'), ('refuse', 'VBP'),
 ('to', 'TO'), ('permit', 'VB'), ('us', 'PRP'),
 ('to', 'TO'), ('obtain', 'VB'), ('the', 'DT'),
 ('refuse', 'NN'), ('permit', 'NN')]
```

Representing tagged tokens

NLTK uses tuples to represent word, tag pairs:

```
>>> tagged_token = nltk.tag.str2tuple('fly/NN')
>>> tagged_token
('fly', 'NN')
>>>
>>> sent = 'They/PRP refuse/VBP to/TO permit/VB
           us/PRP to/TO obtain/VB the/DT
           refuse/NN permit/NN'
>>> [nltk.tag.str2tuple(t) for t in sent.split()]
[('They', 'PRP'), ('refuse', 'VBP'), ('to', 'TO'),
 ('permit', 'VB'), ('us', 'PRP'), ('to', 'TO'),
 ('obtain', 'VB'), ('the', 'DT'), ('refuse', 'NN'),
 ('permit', 'NN')]
```

Python basics

NLTK

Texts

Lists

Distributions

Control structures

Nested Blocks

New data

POS Tagging

Basic tagging

Tagged corpora

Automatic tagging

Reading tagged corpora

NLTK has a variety of corpora to work with (see <http://nltk.org/book/ch02.html>)

```
>>> nltk.corpus.brown.tagged_words()  
[('The', 'AT'), ('Fulton', 'NP-TL'), ...]  
>>> nltk.corpus.brown.tagged_words(simplify_tags=True)  
[('The', 'DET'), ('Fulton', 'NP'), ('County', 'N'), ...]
```

Ways to access information for tagged corpora:

- ▶ `.words()`
[list of words]
- ▶ `.tagged_words()`
[list of (word,tag) pairs]
- ▶ `.sentences()`
[list of list of words]
- ▶ `.tagged_sentences()`
[list of list of (word,tag) pairs]
- ▶ `.paras()`
[list of list of list of words]
- ▶ `.tagged_paras()`
[[list of list of list of (word,tag) pairs]

Python basics

NLTK

Texts

Lists

Distributions

Control structures

Nested Blocks

New data

POS Tagging

Basic tagging

Tagged corpora

Automatic tagging

Automatic POS tagging

Most Frequent Tag Tagger

```
nlk.DefaultTagger():
```

```
>>> raw = 'I do not like green eggs and ham, I \
          do not like them Sam I am!'
>>> tokens = nltk.word_tokenize(raw)
>>> default_tagger = nltk.DefaultTagger('NN')
>>> default_tagger.tag(tokens)
[('I', 'NN'), ('do', 'NN'), ('not', 'NN'), ...]
```

For stored data (lists of lists of word/tag pairs), you can use `.evaluate()`

```
>>> brown_tagged_sents =
      brown.tagged_sents(categories='news')
>>> default_tagger.evaluate(brown_tagged_sents)
0.13089484257215028
```

Python basics

NLTK

Texts

Lists

Distributions

Control structures

Nested Blocks

New data

POS Tagging

Basic tagging

Tagged corpora

Automatic tagging

Automatic POS tagging

Regular Expression Tagger

Regular expressions capture patterns compactly

```
patterns = [  
    ...     (r'.*ing$', 'VBG'),           # gerunds  
    ...     (r'.*ed$', 'VBD'),           # simple past  
    ...     (r'.*es$', 'VBZ'),           # 3rd sg. pres.  
    ...     (r'.*ould$', 'MD'),          # modals  
    ...     (r'.*\'s$', 'NN$'),          # possessive nouns  
    ...     (r'.*s$', 'NNS'),            # plural nouns  
    ...     (r'^-?[0-9]+(.[0-9]+)?$', 'CD'), # cardinal #s  
    ...     (r'.*', 'NN')                # nouns (default)  
    ... ]  
  
>>> regexp_tagger = nltk.RegexpTagger(patterns)
```

Note that the patterns are applied *in order*

Automatic POS tagging

Regular Expression Tagger (2)

```
>>> brown_sents = brown.sents(categories='news')
>>> regexp_tagger.tag(brown_sents[3])
[(''', 'NN'), ... ('such', 'NN'),
 ('reports', 'NNS'), ... ('considering', 'VBG'),
 ('the', 'NN'), ...]
>>>
>>> regexp_tagger.evaluate(brown_tagged_sents)
0.20326391789486245
```

N-gram tagging

Unigram tagging

`nltk.UnigramTagger` learns the most frequent tag for every word:

```
>>> size = int(len(brown_tagged_sents) * 0.9)
>>> size
4160
>>> train_sents = brown_tagged_sents[:size]
>>> test_sents = brown_tagged_sents[size:]
>>> unigram_tagger = nltk.UnigramTagger(train_sents)
>>> unigram_tagger.evaluate(test_sents)
0.8110236220472441
```

Python basics

NLTK

Texts

Lists

Distributions

Control structures

Nested blocks

New data

POS Tagging

Basic tagging

Tagged corpora

Automatic tagging

N-gram tagging

Bigram tagging

`nltk.BigramTagger` learns the most frequent tag for every bigram:

```
>>> bigram_tagger = nltk.BigramTagger(train_sents)
>>> bigram_tagger.tag(brown_sents[2007])
[('Various', 'JJ'), ('of', 'IN'), ('the', 'AT'), ...]
>>>
>>> bigram_tagger.evaluate(test_sents)
0.10216286255357321
```

Note that bigrams which are unseen are assigned nothing

N-gram tagging

Combining taggers

Use the best information if you have it:

```
>>> t0 = nltk.DefaultTagger('NN')
>>> t1 = nltk.UnigramTagger(train_sents, backoff=t0)
>>> t2 = nltk.BigramTagger(train_sents, backoff=t1)
>>> t2.evaluate(test_sents)
0.8447124489185687
```

Unknown words can (also) be handled via regular expressions and be better integrated into contextual information